



Training pass 1

rate = 0.5 — half the hidden units are off



Training pass 2

rate = 0.5 — half the hidden units are off



Same network, same weights — different units switched off each pass.

So no unit can rely on any particular other one being there. That is the whole mechanism.

At test time nothing is dropped; the activations were already scaled up by $1/(1-\text{rate})$ during training.