

1 · Intermediate activations

how successive layers
transform an input



2 · Filter visualisation

what pattern each
filter responds to



3 · Class activation maps

which part of the image
drove the decision



4 · Latent space

which inputs the model
finds semantically similar