

Max pooling — fine for classification

Take a 2×2 window

Return one value

**Which of the four positions
it came from is destroyed**

Strided convolution — needed for segmentation

Slide the kernel
with stride 2

Downsample by learning

**Location information
is retained**