


detection



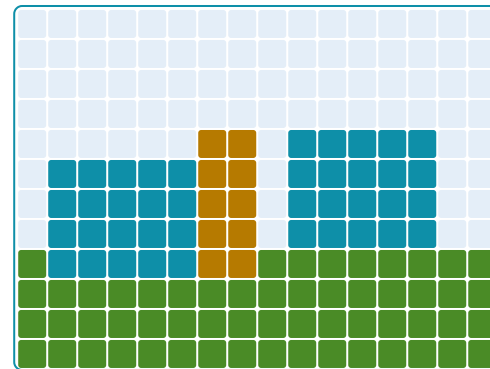
3 objects $\times$ (x, y, w, h, class)

15 numbers

17 476×

more output
for the same
question

segmentation



512 $\times$ 512 pixels, one label each

262 144 labels

Both answer the same question. One of them is what the last layer has to emit — and what a person has to annotate, by hand, for every training image.