

Unsupervised pretraining

guess the next token

The Transformer

from the 2017 paper

A lot of text

thousands of
self-published books

GPT

Generative Pretrained
Transformer

State of the art across
a wide array of tasks
with no per-task
architecture changes

