

1. Supervised fine-tuning

handwritten prompts
and responses



2. Rank responses

human evaluators order them
most to least helpful



3. Reward model

a smaller Transformer:
sequence in, one float out



4. Reinforcement learning

generate, score, update



5. Iterate

new responses, new
rankings

