

1. Encode

tokenize the prompt,
embed it

2. Denoise

pure noise to latents,
conditioned on the text

3. Decode and scale

$[-1, 1]$ back to $[0, 255]$