

Prompt

a short description
of the image



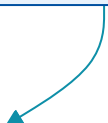
Pretrained text encoder

a Transformer encoder,
like RoBERTa



Text embeddings

Noisy image



Denoising model

two inputs now



Predicted noise mask