

Two ways to spread work
across devices

```
graph TD; A[Two ways to spread work across devices] --> B[Data parallelism]; A --> C[Model parallelism]; B --> D[Speed. Requires the model to fit on one device.]; C --> E[Size. Used when the model does not fit anywhere.];
```

Data parallelism

one model replicated on
every device; each processes
a different sub-batch

Speed. Requires the model
to **fit on one device**.

Model parallelism

different parts of one model
on different devices, working
on the same batch

Size. Used when the model
does not fit anywhere.