

1. Officer photographs the application in the field

2. Agent pulls customer and transaction data
REST API, via MCP

3. Agent runs the analysis
SQL over 12 months

4. Agent calls the credit model
classical ML, returns a score and the reason codes

5. Agent retrieves the policy
which rules apply to this case

6. Agent drafts a recommendation
with every figure cited

7. HUMAN decides
the agent never approves credit

8. Trace written to the audit log