



Institut Teknologi Bandung

Directorate of Continuing Professional Education

in partnership with

CHAPTER 2

The Mathematical Building Blocks of Neural Networks

Tensors, tensor operations, and gradient-based optimisation — explained through code you can run rather than notation you have to decode.

Duration 3 hours (2 sessions)

Instructor Rahman Indra Kesuma, S.Kom., M.Cs. · Prof. Bambang Riyanto Trilaksono

Source Chollet & Watson, Deep Learning with Python 3e — chapter 2

Session Objectives

By the end of this session, participants can:

- 1 Run the first MNIST example end to end and name the role of **layers, loss, optimizer, and metrics** in each line of it.
- 2 Read a tensor's **rank, shape, and dtype**, and map real data — vectors, timeseries, images, video — onto its tensor shape.
- 3 Explain **element-wise operations, broadcasting, the tensor product, and reshaping**, together with what each one means geometrically.
- 4 Describe **derivatives, gradients, mini-batch SGD, and the chain rule**, and point to where each lives in a computation graph.
- 5 Rewrite MNIST **from scratch** — layer, model, batch generator, and training loop — without calling `fit()`.

BOOK

[Chapter 2 — full text](#)

NOTEBOOK

[01_first_mnist.ipynb](#)

NOTEBOOK

[02_tensors_and_operations.ipynb](#)

NOTEBOOK

[03_gradients_and_sgd.ipynb](#)

NOTEBOOK

[04_mnist_from_scratch.ipynb](#)

REPO

[Author's official notebook](#)

One example, taken apart all the way down

The chapter travels a full circle: run MNIST with `fit()`, dismantle every piece of it, then **rewrite the whole thing from scratch** and show the result holds up.

1 · First example

MNIST in ten lines. Working first, understood second.

2.1

2 · Tensors

Rank, shape, dtype, slicing, the batch axis, real-world data.

2.2

3 · Tensor operations

Element-wise, broadcasting, matmul, reshape — and their geometry.

2.3

4 · The engine

Derivatives, gradients, SGD, the chain rule, autodiff.

2.4–2.6

Why this chapter shows code instead of notation

Runnable code is the most precise, unambiguous description of a mathematical operation.

The working principle of chapter 2

Every concept here is introduced twice: once as an idea, and once as a few lines you can execute and inspect. If the two ever disagree, **the code is right**

01

A first look at a neural network

MNIST — the "hello world" of deep learning.

The data: 60,000 to train on, 10,000 to be judged on

listing 2.1 – loading MNIST

PYTHON

```
from keras.datasets import mnist

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

print(train_images.shape, train_images.dtype)
print(len(train_labels), train_labels[:10])
print(test_images.shape)
```

OUTPUT

```
(60000, 28, 28) uint8
60000 [5 0 4 1 9 2 1 3 1 4]
(10000, 28, 28)
```

Three words used precisely from here on

TERM	WHAT IT MEANS HERE
Sample	One data point — a single 28×28 image.
Class	One category — a digit from 0 to 9.
Label	The class attached to one particular sample.

These are used consistently for the rest of the book, and mixing them up is the fastest way to misread an error message later.

The model: two layers, and what a layer is for

listing 2.2 — the model

PYTHON

```
import keras
from keras import layers

model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax"),
])
```

The final `softmax` layer emits **10 probability scores that sum to 1** — one per digit class. Note there is no `input_shape` anywhere; Keras infers it on the first call, which chapter 3 explains.

Compilation: three choices, and only three

listing 2.3 – compile

PYTHON

```
model.compile(  
    optimizer="adam",  
    loss="sparse_categorical_crossentropy",  
    metrics=["accuracy"],  
)
```

Loss

The feedback signal the network steers by. **This is what gets minimised.**

Optimizer

The mechanism by which the network updates itself from that signal.

Metrics

Watched during training, but **never optimised for.**

Preprocessing: flatten, cast, and scale

listing 2.4 – preparing the data

PYTHON

```
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255

test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255


print(train_images.shape, train_images.dtype, train_images.min(), train_images.max())
```

OUTPUT

```
(60000, 784) float32 0.0 1.0
```

Two things changed: the shape went from `(60000, 28, 28)` to `(60000, 784)`, and the values went from `0..255` integers to `0..1 floats`. Chapter 6 explains why the second matters so much.

The whole first example, as a shape pipeline



Every arrow changes either the shape or the dtype. Nothing else happens in the first example.

Training, and the first crack to notice

listing 2.5 – fit

PYTHON

```
model.fit(train_images, train_labels, epochs=5, batch_size=128)

test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"test accuracy: {test_acc:.3f}")
```

OUTPUT

```
Epoch 1/5
469/469 ---- 3s 5ms/step - accuracy: 0.8747 - loss: 0.4358
Epoch 5/5
469/469 ---- 2s 5ms/step - accuracy: 0.9890 - loss: 0.0361

313/313 ---- 1s 2ms/step - accuracy: 0.9780 - loss: 0.0745
test accuracy: 0.978
```

98.9% on data it trained on, **97.8%** on data it had never seen. That gap is not noise — it is **overfitting**, and chapter 5 treats it as the central problem of the field.

Making a prediction, and reading it

listing 2.6 – predicting

PYTHON

```
test_digits = test_images[0:10]
predictions = model.predict(test_digits)

print(predictions[0].argmax())    # which class?
print(predictions[0].max())      # how confident?
print(test_labels[0])           # what was the truth?
```

OUTPUT

```
7
0.99993
7
```

Right answer, and near-total confidence. Chapter 5 will show why **confidence and correctness are not the same thing**

02

Data representations: tensors

A container for data. Three attributes, and one special axis.

Rank, shape, dtype

Rank 0 — scalar

shape `()`

7

a single number

Rank 1 — vector

shape `(4,)`

7 2 9 4

a row of numbers

Rank 2 — matrix

shape `(3, 4)`

7	2	9	4
1	8	3	6
5	0	2	7

rows of vectors

Rank 3 — tensor

shape `(2, 3, 4)`

7	2	9	4
1	8	3	6
5	0	2	7

a stack of matrices

Rank is how many indices it takes to reach one number: `t`, `t[i]`, `t[i][j]`, `t[i][j][k]`.

The same four numbers 7, 2, 9, 4 appear at every rank. What changes is how many indices it takes to reach one of them.

Reading the drawing: what each rank actually is

RANK	NAME	SHAPE	TO REACH ONE NUMBER	EXAMPLE
0	scalar	<code>()</code>	<code>t</code>	a loss value — one number per batch
1	vector	<code>(4,)</code>	<code>t[i]</code>	one MNIST image after flattening is <code>(784,)</code>
2	matrix	<code>(3, 4)</code>	<code>t[i][j]</code>	a batch of flattened images, <code>(128, 784)</code>
3	tensor	<code>(2, 3, 4)</code>	<code>t[i][j][k]</code>	a batch of images before flattening, <code>(128, 28, 28)</code>

Rank is a count of axes, not a size. A `(60000, 28, 28)` tensor and a `(2, 3, 4)` tensor are both rank 3, and every rule you learn about one applies to the other. **That is the whole reason the abstraction is worth having.**

The word *dimension* is used for two different things and causes real confusion: a rank-3 tensor has **3 axes**, while a `(784,)` vector is often called *784-dimensional*. Say **rank** for axes and **shape** for sizes, and the ambiguity disappears.

Reading them off a real tensor

the three attributes

PYTHON

```
(train_images, train_labels), _ = mnist.load_data()

print(train_images.ndim)    # rank: how many axes
print(train_images.shape)   # shape: how long each axis is
print(train_images.dtype)   # dtype: what the entries are
```

OUTPUT

```
3
(60000, 28, 28)
uint8
```

The terminology trap that costs people an afternoon

A **5-dimensional vector** is not a **5-dimensional tensor**. The first has **one axis holding five numbers**; the second has **five axes**

Misreading this makes shape error messages unintelligible, and shape errors are the most common failure you will hit in the exercises.

Slicing tensors

listing 2.7-2.8 – slices

PYTHON

```
print(train_images[10:100].shape)      # 90 images
print(train_images[:, 14:, 14:].shape)  # bottom-right 14x14 corner
print(train_images[:, 7:-7, 7:-7].shape) # centre 14x14, via negative indices
```

OUTPUT

```
(90, 28, 28)
(60000, 14, 14)
(60000, 14, 14)
```

The batch axis is always axis 0

listing 2.9 – batches

PYTHON

```
batch = train_images[:128]      # batch 0
batch = train_images[128:256]    # batch 1

n = 3
batch = train_images[128 * n : 128 * (n + 1)]    # batch n
print(batch.shape)
```

OUTPUT

```
(128, 28, 28)
```

Axis 0 is the **samples axis** — and because of batching it is also called the **batch axis**. **Every shape error you will read starts with identifying this axis.**

Real-world data, and the shape it takes

KIND OF DATA	RANK	SHAPE	EXAMPLE FROM THE BOOK
Vector data	2	(samples, features)	100,000 people $\times$ (age, gender, income) $\rightarrow$ (100000, 3)
Timeseries	3	(samples, timesteps, features)	250 days $\times$ 390 minutes $\times$ 3 values $\rightarrow$ (250, 390, 3)
Images	4	(samples, h, w, channels)	128 RGB images at $256 \times 256 \rightarrow$ (128, 256, 256, 3)
Video	5	(samples, frames, h, w, channels)	4 clips $\times$ 240 frames $\times$ $144 \times 256 \times$ RGB $\rightarrow$ (4, 240, 144, 256, 3)

Channels - last or channels - first — and why it bites

- ♦ **Channels-last** `(..., h, w, c)` — the TensorFlow and JAX convention.
- ♦ **Channels-first** `(..., c, h, w)` — the PyTorch convention.

Keras 3 picks between them with `image_data_format`.

Getting this wrong produces shape errors that look nonsensical, because both orderings are *valid* — just not for the same layer.

425 MB

its size at float32

106,168,320

values in that 60-second video example

03

The gears: tensor operations

It all reduces to a handful of operations — each with a geometric meaning.

A Dense layer is three operations, and nothing else

the whole of a Dense layer

PYTHON

```
output = relu(matmul(input, W) + b)
#           ^         ^         ^
#           |         |         +-- addition (with broadcasting)
#           |         +----- tensor product
#           +----- element-wise operation
```

`relu(x)` is simply `max(x, 0)`: it zeroes out anything negative and leaves the rest alone.

Element-wise operations, written the slow way

listing 2.10 – naive relu

PYTHON

```
def naive_relu(x):  
    assert len(x.shape) == 2  
    x = x.copy()  
    for i in range(x.shape[0]):  
        for j in range(x.shape[1]):  
            x[i, j] = max(x[i, j], 0)  
    return x
```

Correct, readable — and unusably slow, as the next slide shows.

...and the fast way, which is the same operation

the vectorised version

PYTHON

```
z = x + y  
z = np.maximum(z, 0.0)
```

0.02 s


vectorised NumPy, 1,000 iterations

2.45 s

naive Python loops, 1,000 iterations

Roughly **100×**, and it is not a language effect. NumPy hands the work to BLAS routines written in C and Fortran; on a GPU the same operation runs as **fully vectorised CUDA**

Broadcasting: fitting a small shape to a big one



Two steps: add axes until the ranks match, then repeat along the new axes.

No memory is actually duplicated — the repetition is **algorithmic, not physical**

The rule, and where it quietly goes wrong

listing 2.11 – broadcasting in anger

PYTHON

```
X = np.random.random((64, 3, 32, 10))
y = np.random.random((32, 10))

z = np.maximum(X, y)      # y is broadcast across the first two axes
print(z.shape)
```

OUTPUT

```
(64, 3, 32, 10)
```

This is the most common source of **silent** bugs: the shapes line up, the code runs, and the **meaning is wrong**. When a result looks strange, print the shapes first.

The tensor product, and its one compatibility rule

matmul

PYTHON

```
z = np.matmul(x, y)
z = x @ y           # the same thing, shorter

# compatibility:  x.shape[1] == y.shape[0]
# result shape:  (x.shape[0], y.shape[1])

# (a, b, c, d) @ (d,)    -> (a, b, c)
# (a, b, c, d) @ (d, e)  -> (a, b, c, e)
```

Visually: line the two up as rectangles — **the width of the first must match the height of the second**

Reshaping rearranges; it never invents or destroys

listing 2.12 – reshape and transpose

PYTHON

```
x = np.array([[0., 1.],
              [2., 3.],
              [4., 5.]])      # shape (3, 2)

print(np.reshape(x, (6,)).shape)    # flattened
print(np.reshape(x, (2, 3)).shape)  # regrouped


print(np.transpose(np.zeros((300, 20))).shape)  # rows and columns exchanged
```

OUTPUT

```
(6,)
(2, 3)
(20, 300)
```

This is exactly what `train_images.reshape((60000, 784))` did in the first example: **the pixels never changed, only their arrangement**

Every tensor operation is a geometric one



Dashed: before. Solid: after. Every layer in a network does the last one — a matrix, then a vector.

One shape, four operations. The dashed outline is where it started.

Reading the drawing: what moved, and what did it

OPERATION	WHAT IT DOES TO THE SHAPE	WRITTEN AS
Translation	Slides it. Every point moves by the same vector; nothing turns or stretches.	$y = x + b$
Rotation	Turns it about the origin. Lengths and angles survive — the shape is the same shape, pointing elsewhere.	$y = R @ x$
Scaling	Stretches each axis by its own factor. Here x by 1.5 and y by 0.6, which is why it comes out squat.	$y = D @ x$
Affine	Any matrix, then a vector. Straight lines stay straight and parallel lines stay parallel — but angles do not survive.	$y = W @ x + b$

A Dense layer is the fourth row. $W @ x + b$ is an affine transform, and nothing else. A stack of them with no activation between is still one affine transform — which is the next slide, and the reason activations exist.

Why an activation function is not optional

two affine layers collapse into one

PYTHON

```
affine2(affine1(x)) == (W2 @ W1) @ x + (W2 @ b1 + b2)
#                      \-----/
#                      one matrix
#                      \-----/
#                      one vector
```

So a deep stack of `Dense` layers with no activation is **secretly a single linear model**, however many layers you give it. Nonlinearities such as `relu` are what make the hypothesis space rich.

Deep learning as uncrumpling a paper ball

It happens through a long series of **small, elementary geometric moves** — like uncrumpling a paper ball with successive finger movements. **Each layer disentangles the data a little.**


Chapter 5 makes this precise, under the name of the *manifold hypothesis*.

04

The engine: gradient-based optimisation

Derivatives, gradients, SGD, and the chain rule.

The training loop, and which step is hard



Four steps, repeated over batches. Only the fourth is difficult.

Tuning each coefficient by hand is impossible — modern networks have **millions to billions of parameters**. Gradient descent solves for all of them at once.

A derivative is the slope of a local approximation

- ♦ a negative $\rightarrow$ increasing x **decreases** $f(x)$.
- ♦ a positive $\rightarrow$ increasing x **increases** $f(x)$.
- ♦ Its magnitude says how fast the change happens.

So to make f smaller, move x **in the direction opposite its derivative**. That single sentence is the whole of optimisation.

A gradient is a derivative for tensors

what the gradient is a gradient of

PYTHON

```
y_pred = matmul(x, W)
loss_value = loss(y_pred, y_true)

# For fixed x and y_true this is just loss_value = f(W)
# grad(loss_value, W0) is a tensor SHAPED LIKE W, whose every coefficient says
# in which direction, and how strongly, the loss moves if you nudge that weight.
```

To lower the loss, step against the gradient: $W_1 = W_0 - \text{step} * \text{grad}(f(W_0), W_0)$

Why this is done by iteration and not by solving

Because for a network with millions of parameters that is **analytically intractable**. Iteration is not a shortcut; it is the only available route.

Mini-batch stochastic gradient descent

- 1 Draw a **random** batch of samples and targets. (The word *stochastic* comes from this randomness.)
- 2 Forward pass → predictions.
- 3 Compute the loss.
- 4 **Backward pass** → gradient of the loss with respect to the parameters.
- 5 `W -= learning_rate * gradient`.

Repeat. Each step lowers the loss a little; enough steps and the model fits.

Learning rate: the one number that decides whether it works

Too small

Convergence crawls — many steps for very little progress. It can look as though training has stalled.

Momentum


Update using the current gradient **and** previous updates — like a ball rolling downhill with enough speed to cross a shallow dip.

Too large

Updates overshoot wildly; the loss jumps around and **never comes down**.

The variants that improve on plain SGD are called **optimizers**: SGD with momentum, Adagrad, RMSprop, Adam.

The chain rule, applied backwards over a graph



Solid arrows are the forward pass; dashed arrows are the backward pass, each carrying one local derivative.

Multiply the local derivatives along a path and you have the gradient. Applying the chain rule this way over a network is backpropagation.

Two details that make the graph work

- ♦ When there are **several paths** from node **a** to node **b**, the contributions of the paths are **summed**.
- ♦ A **computation graph** is a directed acyclic graph of operations. It makes computation itself into data — a structure a machine can read and manipulate.

Modern frameworks walk that graph for you. That is **automatic differentiation**, and it is why **you will never write backpropagation by hand**

Autodiff in three lines

listing 2.19 – a first derivative

PYTHON

```
import tensorflow as tf

x = tf.Variable(3.0)
with tf.GradientTape() as tape:
    y = 2 * x + 3

print(tape.gradient(y, x))    # dy/dx
```

OUTPUT

```
tf.Tensor(2.0, shape=(), dtype=float32)
```

The answer is 2, which is what you would get by hand — except nothing here was told the rule for differentiating a line.

...and a second derivative, which is a physics answer

listing 2.20 – nested tapes

PYTHON

```
time = tf.Variable(0.0)

with tf.GradientTape() as outer_tape:
    with tf.GradientTape() as inner_tape:
        position = 4.9 * time ** 2
        speed = inner_tape.gradient(position, time)

    acceleration = outer_tape.gradient(speed, time)
    print(acceleration)
```

OUTPUT

```
tf.Tensor(9.8, shape=(), dtype=float32)
```


`position = 4.9 · t2` is free fall. Its second derivative is the **acceleration due to gravity**, and autodiff returned 9.8 without ever being given the formula.

05

Rewriting the first example from scratch

No `fit()`, no built-in `Dense`. The result has to hold up.

What we are about to rebuild



Four small classes and one function replace Sequential, Dense, and fit() between them.

A layer, written out

listing 2.21 – NaiveDense

PYTHON

```
import keras
from keras import ops

class NaiveDense:
    def __init__(self, input_size, output_size, activation=None):
        self.activation = activation
        self.W = keras.Variable(shape=(input_size, output_size), initializer="uniform")
        self.b = keras.Variable(shape=(output_size,), initializer="zeros")

    def __call__(self, inputs):
        x = ops.matmul(inputs, self.W) + self.b
        return self.activation(x) if self.activation is not None else x

    @property
    def weights(self):
        return [self.W, self.b]
```

W starts random and **b** starts at zero — exactly the starting condition chapter 1 described.

A model is just layers, applied in order

listing 2.22 - NaiveSequential

PYTHON

```
class NaiveSequential:
    def __init__(self, layers):
        self.layers = layers

    def __call__(self, inputs):
        x = inputs
        for layer in self.layers:
            x = layer(x)
        return x

    @property
    def weights(self):
        return [w for layer in self.layers for w in layer.weights]
```

Note how little there is to it: a model is **a list of layers and a for loop**. Everything else in Keras is convenience on top.

One training step is the four figures, in code

listing 2.24 – one_training_step

PYTHON

```
import tensorflow as tf
from keras import optimizers

optimizer = optimizers.SGD(learning_rate=1e-3)

def one_training_step(model, images_batch, labels_batch):
    with tf.GradientTape() as tape:
        predictions = model(images_batch)           # forward
        loss = ops.sparse_categorical_crossentropy(labels_batch, predictions)
        average_loss = ops.mean(loss)               # loss
    gradients = tape.gradient(average_loss, model.weights) # gradients
    optimizer.apply_gradients(zip(gradients, model.weights)) # update
    return average_loss
```

Four lines, four ideas: **forward**, **loss**, **gradients**, **update**. Every training loop in the rest of the book is a variation on this shape.

And the loop around it

listing 2.26 – fit, by hand

PYTHON

```
def fit(model, images, labels, epochs, batch_size=128):
    for epoch in range(epochs):
        print(f"Epoch {epoch}")
        gen = BatchGenerator(images, labels, batch_size)
        for i in range(gen.num_batches):
            images_batch, labels_batch = gen.next()
            loss = one_training_step(model, images_batch, labels_batch)
            if i % 100 == 0:
                print(f"  loss at batch {i}: {loss:.2f}")
```

`BatchGenerator` is a dozen lines that hand out consecutive slices — the batch axis from section 2.2, put to work.

Does it work? Run it and see

listing 2.27 – evaluating it

PYTHON

```
model = NaiveSequential([
    NaiveDense(input_size=28 * 28, output_size=512, activation=ops.relu),
    NaiveDense(input_size=512, output_size=10, activation=ops.softmax),
])
fit(model, train_images, train_labels, epochs=10, batch_size=128)

predictions = model(test_images)
predicted_labels = ops.argmax(predictions, axis=1)
print(f"accuracy: {ops.mean(predicted_labels == test_labels):.2f}")
```

OUTPUT

```
Epoch 0
  loss at batch 0: 6.19
  loss at batch 400: 2.21
Epoch 9
  loss at batch 0: 0.36
  loss at batch 400: 0.34
accuracy: 0.90
```

90% versus 97.8% — and why that is the lesson

The hand-written version scores **lower**, and that is not a bug. It uses plain SGD at `learning_rate=1e-3`; the first example used **Adam**

So the gap measures **the choice of optimizer**, not the quality of the implementation. It is one of the clearest demonstrations in the book of how much that choice is worth.

What has to survive this chapter

- 1 **Tensors** are containers with **rank**, **shape**, **dtype**. Axis 0 is the samples axis, and therefore the batch axis.
- 2 **Tensor operations are geometric**. A Dense layer without an activation is an affine transform, and stacking those still gives an affine transform.
- 3 **Learning** means finding weight values that minimise the loss on the training data.
- 4 **Mini-batch SGD** updates weights from gradients computed on random batches.
- 5 **Chain rule + autodiff = backpropagation**, carried out over a computation graph.
- 6 **Loss** measures success at the task; **optimizer** decides how the descent is performed — and that choice is worth several points of accuracy.

NOTEBOOK

[01_first_mnist.ipynb](#)

NOTEBOOK

[04_mnist_from_scratch.ipynb](#)

NEXT

[Chapter 3 — The frameworks](#)