



Institut Teknologi Bandung

Directorate of Continuing Professional Education

in partnership with

CHAPTER 10

Interpreting What ConvNets Learn

Four ways to open the box — and the argument that for ConvNets, the box was never as closed as people say.

Duration 2.5 hours

Instructor Prof. Bambang Riyanto Trilaksono

Source Chollet & Watson, Deep Learning with Python 3e — chapter 10

Session Objectives

By the end of this session, participants can:

- 1 Build a **multi-output model** that returns intermediate activations, and read what each depth of layer is doing.
- 2 Explain **information distillation** and the rising sparsity of activations with depth.
- 3 Visualise what a filter responds to using **gradient ascent in input space**, and say how it differs from ordinary training.
- 4 Produce a **Grad-CAM heatmap** and use it to explain — or debug — a particular classification.
- 5 Project a model's **latent space** to two dimensions to find outliers and semantically ambiguous samples.
- 6 Say which technique answers which question, so the right one is reached for in a real investigation.

BOOK

[Chapter 10 — full text](#)

NOTEBOOK

[01_intermediate_activations.ipynb](#)

NOTEBOOK

[02_filter_visualisation.ipynb](#)

NOTEBOOK

[03_grad_cam.ipynb](#)

NOTEBOOK

[04_latent_space.ipynb](#)

REPO

[Author's official notebook](#)

The black-box claim, examined


Partially true for **certain types** of model. **Definitely not true for ConvNets** — their representations are highly amenable to visualisation, largely because they are **representations of visual concepts**

Since 2013 a wide array of techniques has been developed. This chapter covers the four most accessible and useful.

Why this matters beyond curiosity

This is especially relevant where deep learning **complements human expertise** rather than replacing it — medical imaging being the example the book names. **An unexplained disagreement is not usable in that setting.**

The four techniques, and what each answers



Roughly in order of effort. The first uses your own small ConvNet; the next two use a pretrained Xception.

01

Visualising intermediate activations

What each layer does to one particular image.

What we are looking at

Feature maps have three dimensions — width, height, and channels. Because **each channel encodes a relatively independent feature**, the right way to look at them is **one channel at a time, as a 2D image**

A model whose outputs are the layers themselves

listing 10.3 – the activation model

PYTHON

```
import keras

model = keras.models.load_model("convnet_from_scratch_with_augmentation.keras")

layer_outputs, layer_names = [], []
for layer in model.layers:
    if isinstance(layer, (keras.layers.Conv2D, keras.layers.MaxPooling2D)):
        layer_outputs.append(layer.output)
        layer_names.append(layer.name)

activation_model = keras.Model(inputs=model.input, outputs=layer_outputs)
```

This works **only because the model is Functional**. A subclassed model has no graph to reach into — exactly the trade-off chapter 7 described.

Running it, and plotting one channel

listing 10.4–10.5 – activations

PYTHON

```
activations = activation_model.predict(img_tensor)  # a list of nine arrays

first_layer_activation = activations[0]
print(first_layer_activation.shape)

import matplotlib.pyplot as plt
plt.matshow(first_layer_activation[0, :, :, 5], cmap="viridis")  # the sixth channel
```

OUTPUT

```
(1, 178, 178, 32)
```

That channel turns out to be a **diagonal edge detector** — though your own channels will differ, since the specific filters a layer learns are **not deterministic**

Three things visible in the full grid

- ① **The first layer is a collection of edge detectors.** At that stage the activations retain almost all the information in the original picture.
- ② **Higher layers become increasingly abstract** and less visually interpretable — encoding concepts like *cat ear* and *cat eye*. They carry less about the image's appearance and more about its class.
- ③ **Sparsity increases with depth.** In the first layer every filter fires; further up, more and more are blank — meaning **that filter's pattern is not present in this image.**

A deep network is an information distillation pipeline



Irrelevant information is filtered out; useful information is magnified and refined.

This is a **universal characteristic** of the representations deep networks learn, not something particular to this model.

The bicycle test

Try drawing a generic bicycle from memory right now. You have seen thousands, and you will almost certainly **not get it remotely right**. The effect is real, and it is the same distillation: your brain abstracted the input and discarded the detail.

02

Visualising ConvNet filters

Gradient ascent in input space.

The idea: train the image instead of the weights



Same machinery, opposite target — and the sign is flipped.

Start from a blank (noisy) input and apply gradient descent to the **image**, maximising the response of one chosen filter. The result is the input that filter likes best.

Setting up: a pretrained model, and its layer names

listing 10.7–10.8 – model and layer names

PYTHON

```
import keras_hub

model = keras_hub.models.Backbone.from_preset("xception_41_imagenet")
preprocessor = keras_hub.layers.ImageConverter.from_preset(
    "xception_41_imagenet", image_size=(180, 180))

for layer in model.layers:
    if isinstance(layer, (keras.layers.Conv2D, keras.layers.SeparableConv2D)):
        print(layer.name)
```

OUTPUT

```
block2_sepconv1
block3_sepconv1
block4_sepconv1
...
block14_sepconv2
```

The names follow Xception's block structure from chapter 9, which makes it easy to sample **shallow, middle, and deep** layers for comparison.

The loss: one filter's mean activation

the loss to maximise

PYTHON

```
layer_name = "block3_sepconv1"
layer = model.get_layer(name=layer_name)
feature_extractor = keras.Model(inputs=model.input, outputs=layer.output)

def compute_loss(image, filter_index):
    activation = feature_extractor(image)
    filter_activation = activation[:, 2:-2, 2:-2, filter_index]  # trim the border
    return ops.mean(filter_activation)
```

Note what this loss is a function of: **the image**. That is the whole trick.

One ascent step, in TensorFlow

listing 10.11 — gradient ascent step

PYTHON

```
@tf.function
def gradient_ascent_step(image, filter_index, learning_rate):
    with tf.GradientTape() as tape:
        tape.watch(image)                # the image is not a Variable, so watch it
        loss = compute_loss(image, filter_index)
        grads = tape.gradient(loss, image) # gradient w.r.t. the IMAGE
        grads = ops.normalize(grads)       # the "gradient normalization trick"
        image += learning_rate * grads     # PLUS: we are ascending, not descending
    return image
```

`tape.watch(image)` is required because only `Variable`s are watched automatically — exactly the rule chapter 3 gave. And the `+=` is what makes it **ascent rather than descent**

The same step in PyTorch and JAX

BACKEND	HOW THE GRADIENT IS OBTAINED
TensorFlow	<code>tape.watch(image)</code> , then <code>tape.gradient(loss, image)</code> .
PyTorch	<code>loss.backward()</code> , then read <code>image.grad</code> .
JAX	Transform the loss function with <code>jax.grad</code> , differentiating with respect to the image argument.

Whichever backend you use, **the image is the thing being differentiated**, and that is the only unusual part.

The loop that produces one filter's pattern

listing 10.14 — generating a pattern

PYTHON

```
img_width = img_height = 200

def generate_filter_pattern(filter_index):
    iterations = 30
    learning_rate = 10.0
    image = keras.random.uniform(
        minval=0.4, maxval=0.6, shape=(1, img_width, img_height, 3))
    for i in range(iterations):
        image = gradient_ascent_step(image, filter_index, learning_rate)
    return image[0]
```

A learning rate of **10.0** would be absurd for training weights. Here it is fine, because **we are not trying to converge to a minimum** — only to move the image somewhere the filter likes.

...and turning the result back into an image

listing 10.15 — deprocess_image

PYTHON

```
def deprocess_image(image):  
    image -= ops.mean(image)  
    image /= ops.std(image)  
    image *= 64  
    image += 128  
    image = ops.clip(image, 0, 255).astype("uint8")  
    return image[25:-25, 25:-25, :]    # crop the border artefacts
```

The final crop removes the edge artefacts that gradient ascent tends to accumulate at the borders.

What the filter banks look like, layer by layer

DEPTH	EXAMPLE LAYER	WHAT THE FILTERS ENCODE
Early	<code>block2_sepconv1</code>	Simple directional edges and colours — sometimes coloured edges.
Middle	<code>block4_sepconv1</code>	Simple textures made from combinations of edges and colours.
Deep	<code>block8_sepconv1</code>	Textures found in natural images: feathers, eyes, leaves .

Chollet's analogy: each layer learns a **filter bank** such that its inputs can be expressed as a combination of those filters — **similar to how a Fourier transform decomposes a signal onto cosines**

03

Class activation heatmaps

Which part of this picture made the model say that?

What a CAM is, and what it is for

Debugging a decision

Particularly a **misclassification** — the field is called model interpretability.


Locating an object

The heatmap tells you **where** in the frame the evidence was.

For a dogs-vs-cats model you could generate one heatmap

for *cat* and another for *dog* over the same picture.

Grad-CAM, in one sentence



Weight how strongly each channel fires here by how much this channel matters to the class.

The result is a spatial map of **how intensely the input image activates the class**. The implementation follows Selvaraju et al.

The worked example: two African elephants

listing 10.17 — preparing the image

PYTHON

```
img_path = keras.utils.get_file(
    fname="elephant.jpg",
    origin="https://img-datasets.s3.amazonaws.com/elephant.jpg",
)

img = keras.utils.load_img(img_path)          # a PIL image
img_array = np.expand_dims(img, axis=0)        # to NumPy, plus a batch axis
```

Everything that follows asks the same question of this picture: **which pixels made it say *African elephant*?**

Getting the gradient of the top class

the TensorFlow version

PYTHON

```
with tf.GradientTape() as tape:
    last_conv_layer_output = last_conv_layer_model(preprocessed_image)
    tape.watch(last_conv_layer_output)
    preds = classifier_model(last_conv_layer_output)
    top_pred_index = ops.argmax(preds[0])
    top_class_channel = preds[:, top_pred_index]

grads = tape.gradient(top_class_channel, last_conv_layer_output)
```

PyTorch and JAX versions differ only in the gradient idiom, exactly as in the filter visualisation.

Pooling and weighting to get the heatmap

listing 10.23 — building the heatmap

PYTHON

```
# one number per channel: how important that channel is to the top class
pooled_grads = np.mean(grads, axis=(0, 1, 2))

last_conv_layer_output = last_conv_layer_output[0].copy()
for i in range(pooled_grads.shape[-1]):
    last_conv_layer_output[:, :, i] *= pooled_grads[i]    # weight each channel

heatmap = np.mean(last_conv_layer_output, axis=-1)        # average across channels
```

Three lines of arithmetic and no new machinery — Grad-CAM is **a weighted average, not a new model**

Normalising and superimposing it

listing 10.24–10.25 — display

PYTHON

```
heatmap = np.maximum(heatmap, 0)
heatmap /= np.max(heatmap)

heatmap = np.uint8(255 * heatmap)
jet_colors = cm.get_cmap("jet")(np.arange(256))[:, :3]
jet_heatmap = keras.utils.array_to_img(jet_colors[heatmap])
jet_heatmap = jet_heatmap.resize((img.shape[1], img.shape[0]))
jet_heatmap = keras.utils.img_to_array(jet_heatmap)

superimposed_img = jet_heatmap * 0.4 + img
keras.utils.array_to_img(superimposed_img).save("elephant_cam.jpg")
```

The result is the picture with the evidence highlighted — the artefact you would actually show a domain expert.

The two questions it answers

Why did it say African elephant?

Because **these pixels** carried the evidence. If they are the wrong pixels, you have found a spurious correlation.

Where is the elephant?

The heatmap localises the object without ever having been trained to do detection.

The first question is the one that matters in review: a model that is right **for the wrong reasons** passes every accuracy check and **fails the first time the background changes**

04


Visualising the latent space

Which inputs does the model consider similar?

Every model embeds its inputs in a latent space

This is chapter 5's manifold hypothesis, now made visible.

How to look at it



Take the activations of one layer as coordinates, then project them down to two dimensions.

Any layer can be used — **each encodes a different latent space**, and they become progressively more semantically organised with depth. The usual choice is the layer just before the classifier.

What it is actually good for

Finding outliers

Points sitting far from every cluster are usually **mislabelled or corrupt** — the noisy data of chapter 5.

Finding ambiguity

Points between two clusters are the samples the model finds **semantically ambiguous** — chapter 5's ambiguous features.

Improving the dataset

Both findings feed straight back into **cleaning the training set and refining the evaluation set**.

Which makes this the technique that closes the loop: it is an interpretation tool whose output is **a concrete list of samples to go and look at**

Which technique answers which question

QUESTION YOU ARE ASKING	TECHNIQUE
<i>What is each layer doing to my input?</i>	Intermediate activations (§10.1)
<i>What pattern does this filter respond to?</i>	Gradient ascent in input space (§10.2)
<i>Which part of the image drove this decision?</i>	Grad-CAM heatmap (§10.3)
<i>Which samples does my model confuse, and which are junk?</i>	Latent-space projection (§10.4)

What has to survive this chapter

- 1 **ConvNets are not black boxes.** Their representations are visual, and therefore visualisable.
- 2 **A deep network is an information distillation pipeline:** less about appearance, more about class, with rising sparsity at depth.
- 3 **Gradient ascent trains the image, not the weights** — same machinery, opposite target.
- 4 **Grad-CAM is a weighted average**, and it answers *why here?* as well as *where?*
- 5 **Latent-space projection finds the samples to go and inspect** — outliers and ambiguous cases.

NOTEBOOK

[03_grad_cam.ipynb](#)

NEXT

[Chapter 11 — Image segmentation](#)