



Institut Teknologi Bandung

Directorate of Continuing Professional Education

in partnership with

CHAPTER 13

Timeseries Forecasting

Where a common-sense baseline beats two deep learning models in a row — and what that teaches about architecture priors.

Duration 3 hours (2 sessions)

Instructor Prof. Bambang Riyanto Trilaksono

Source Chollet & Watson, Deep Learning with Python 3e — chapter 13

Session Objectives

By the end of this session, participants can:

- 1 Name the four common timeseries tasks and say which are supervised.
- 2 Build a **common-sense baseline** for a forecasting problem and evaluate it properly.
- 3 Explain **why a dense network and a 1D ConvNet fail** on this problem — and what that says about architecture priors.
- 4 Explain what an **RNN** is, and how it differs from a feedforward network.
- 5 Say why **SimpleRNN** is not used in practice, and how **LSTM** solves it.
- 6 Apply **recurrent dropout** correctly, and say why the naive version is harmful.
- 7 **Stack** recurrent layers with `return_sequences`, and know when a **bidirectional** layer helps and when it cannot.

BOOK

[Chapter 13 — full text](#)

NOTEBOOK

[01_jena_data_and_baseline.ipynb](#)

NOTEBOOK

[02_dense_and_conv1d.ipynb](#)

NOTEBOOK

[03_lstm_and_gru.ipynb](#)

NOTEBOOK

[04_dropout_stacking_bidirectional.ipynb](#)

REPO

[Author's official notebook](#)

What a timeseries is, and why it is different

Unlike everything so far, working with timeseries means understanding **the dynamics of a system** — its periodic cycles, how it trends, its regular regime, and **its sudden spikes**

Four things you can do with one

Forecasting
what happens next
demand, revenue, weather

Anomaly detection
something unusual
usually unsupervised:
you cannot train on
anomalies you have not seen

Classification
a label for a whole series
bot or human?

Event detection
a specific expected event
"OK Google", "Hey Alexa"


Forecasting is by far the most common, and is what this chapter covers.

Note the parenthetical on anomaly detection: it is usually **unsupervised**, because you often do not know what kind of anomaly you are looking for — so **you cannot train on examples of it**

The running example

The data is the Jena weather dataset — fourteen meteorological quantities recorded every ten minutes over several years. It is small enough to work with and **genuinely hard enough to be instructive**

The split has to respect time



Chapter 5's arrow-of-time pitfall, made concrete.

The data must **not** be shuffled before splitting: all test data has to be later than all training data. And the normalisation statistics come from **the training portion only**

Normalising fourteen quantities on different scales

listing 13.6 – normalising

PYTHON

```
mean = raw_data[:num_train_samples].mean(axis=0)
raw_data -= mean
std = raw_data[:num_train_samples].std(axis=0)
raw_data /= std
```

Note the slice: `[:num_train_samples]`. The statistics are computed on the **first 210,225 timesteps only** — **the same information-leak rule as chapter 4's housing example**

Windows without copying the data



Consecutive samples share most of their timesteps, so materialising each one would waste enormous memory.

Keras has a utility for exactly this — `timeseries_dataset_from_array()` — which **generates the windows on the fly**, keeping only the original array in memory.

Understanding it on a toy example first

the utility, on integers

PYTHON

```
int_sequence = np.arange(10)

dummy_dataset = keras.utils.timeseries_dataset_from_array(
    data=int_sequence[:-3],      # windows come from here
    targets=int_sequence[3:],    # the target is 3 steps ahead
    sequence_length=3,
    batch_size=2,
)

for inputs, targets in dummy_dataset:
    for i in range(inputs.shape[0]):
        print([int(x) for x in inputs[i]], int(targets[i]))
```

OUTPUT

```
[0, 1, 2] 3
[1, 2, 3] 4
[2, 3, 4] 5
[3, 4, 5] 6
[4, 5, 6] 7
```

...then the real thing

the training dataset

PYTHON

```
sampling_rate = 6          # one sample per hour (data is every 10 minutes)
sequence_length = 120      # five days of hourly readings
delay = sampling_rate * (sequence_length + 24 - 1)

train_dataset = keras.utils.timeseries_dataset_from_array(
    raw_data[:-delay],
    targets=temperature[delay:],
    sampling_rate=sampling_rate,
    sequence_length=sequence_length,
    shuffle=True,
    batch_size=256,
    start_index=0,
    end_index=num_train_samples,
)
```

`start_index` and `end_index` are how the three splits are carved out of one array **without shuffling** — **the chronological order is preserved by construction**

01

The baseline, and two failures

Common sense contains information a model has no access to.

Before any model: what would common sense do?

listing 13.9 – the common-sense baseline

PYTHON

```
def evaluate_naive_method(dataset):
    total_abs_err, samples_seen = 0.0, 0
    for samples, targets in dataset:
        # column 1 is temperature; un-normalise it back to degrees Celsius
        preds = samples[:, -1, 1] * std[1] + mean[1]
        total_abs_err += np.sum(np.abs(preds - targets))
        samples_seen += samples.shape[0]
    return total_abs_err / samples_seen

print(f"Validation MAE: {evaluate_naive_method(val_dataset):.2f}")
print(f"Test MAE: {evaluate_naive_method(test_dataset):.2f}")
```

OUTPUT

```
Validation MAE: 2.44
Test MAE: 2.62
```

That is the number to beat

Assume tomorrow is like today and you are off by **two and a half degrees on average**. Not good enough to launch a forecasting service — but **every model from here on has to beat it to justify itself**

The book makes the general point too: for an unbalanced classification task with 90% class A, *always predict A* scores 90%. **Such elementary baselines can prove surprisingly hard to beat.**

Attempt 1 — a small dense network

listing 13.10 — a densely connected model

PYTHON

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Flatten()(inputs)
x = layers.Dense(16, activation="relu")(x)
outputs = layers.Dense(1)(x)          # no activation: a regression output
model = keras.Model(inputs, outputs)

model.compile(optimizer="adam", loss="mse", metrics=["mae"])
history = model.fit(train_dataset, epochs=10,
                    validation_data=val_dataset, callbacks=callbacks)
```

MSE as the loss, MAE as the metric. MSE is used for training because it is **smooth around zero**, which gradient descent needs; MAE is reported because it is the number a human can interpret.

...and it does not beat the baseline

Some validation losses come close to the no-learning baseline, **but not reliably**. This is exactly why the baseline was worth computing: it turns out to be **not easy to outperform**

As the book puts it: **your common sense contains a lot of valuable information to which a machine learning model has no access.**

Why gradient descent could not find the obvious answer

Because your **hypothesis space** is the space of all two-layer networks with that configuration, and the heuristic is **one model among millions** in it. **Like looking for a needle in a haystack.**

The general limitation: **unless the learning algorithm is hardcoded to look for a specific kind of simple model, it can fail to find a simple solution to a simple problem.**

Attempt 2 — a 1D ConvNet

the 1D convolutional model

PYTHON

```
x = layers.Conv1D(8, 24, activation="relu")(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1)(x)
```

OUTPUT

validation MAE: about 2.9 degrees

Worse than the dense model. Two reasons.

Weather is not translation invariant

There are daily cycles, but **morning data behaves differently from evening or night data**. The invariance holds only at one very specific timescale.

Order matters, a lot

The recent past is far more informative than five days ago — and **max pooling and global average pooling destroy order information**.


A worked demonstration of chapter 9's claim: an architecture encodes assumptions, and **when those assumptions are false the architecture actively hurts**

02

Recurrent neural networks

Treat the data as what it is: a sequence where order matters.

Every network so far had no memory



The dense model removed time; the convolutional model destroyed order. Neither treated the data as a sequence.

The biological analogy the book uses

Biological intelligence processes information **incrementally, while maintaining an internal model** built from past information and constantly updated. An RNN adopts the same principle — **in an extremely simplified version**

An RNN is a network with an internal loop




Figure 13.6 — it iterates through the sequence, keeping a state describing what it has seen so far.

The state is **reset between independent sequences**, so one sequence is still one data point. What changes is that **the data point is no longer processed in a single step**

Why you will rarely use SimpleRNN

The cause is **the vanishing gradient problem** — the same effect chapter 9 described for very deep feedforward networks. Studied by Hochreiter, Schmidhuber and Bengio in the early 1990s.


Keras offers two alternatives designed to fix it: **LSTM** and **GRU**.

LSTM: a conveyor belt running beside the sequence

Imagine a **conveyor belt running parallel to the sequence**. Information can jump onto it at any point, be carried to a later timestep, and jump off **intact** when needed — **preventing older signals from gradually vanishing**

The book points out this should remind you of **residual connections** from chapter 9. It is very nearly the same idea, applied along time rather than along depth.

The carry track, in a picture



Figures 13.8–13.9 — the carry **C** modulates both the next output and the next state.

That is all the structural difference amounts to: a **SimpleRNN** cell, plus **an extra data flow that crosses timesteps without being squashed at every step.**

And it finally beats the baseline

a simple LSTM model

PYTHON

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(16)(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)

model.compile(optimizer="adam", loss="mse", metrics=["mae"])
history = model.fit(train_dataset, epochs=10,
                    validation_data=val_dataset, callbacks=callbacks)
```

2.44

common-sense baseline, validation MAE

2.36

LSTM, validation MAE

The whole chapter, as one ladder



Two deep learning models lose to the heuristic before one beats it.

The margin over the baseline is **small**, and the chapter does not pretend otherwise. The lesson is not that LSTMs are powerful; it is that **the right architecture prior is what made the difference**

03

Getting more out of RNNs

Recurrent dropout, stacking, and bidirectional layers.

Dropout in a recurrent layer is not obvious

It has long been known that applying dropout **before** a recurrent layer **hinders learning rather than helping regularisation**

Yarin Gal's result: hold the mask constant



The same pattern of dropped units at every timestep, rather than a fresh random one.

Determined in 2015 as part of Gal's PhD thesis on Bayesian deep learning. He did the research **using Keras**, and helped build the mechanism directly into its recurrent layers.

The two arguments that implement it

recurrent dropout

PYTHON

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
x = layers.Dropout(0.5)(x)           # regularise the LSTM's OUTPUT too
outputs = layers.Dense(1)(x)
```

ARGUMENT

dropout

recurrent_dropout

WHAT IT DROPS


The layer's **input** units.

The layer's **inner recurrent activations** — a temporally constant mask, which is the part Gal's result is about.

Stacking: when you stop overfitting, add capacity

Recurrent layer stacking is the classic way to do it. Not long ago **Google Translate was powered by a stack of seven large LSTM layers — which is huge**

The one argument that makes stacking work



Intermediate layers must emit a full sequence; only the last one emits a single vector.

a stack of two GRU layers

PYTHON

```
x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs)
x = layers.GRU(32, recurrent_dropout=0.5)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
```

`return_sequences=True` makes a layer return a **rank-3 tensor** — its output at every timestep — instead of only the last one. Omit it on an intermediate layer and **the next layer has no sequence to read**

GRU, and why the book switches to it here

In practice the choice between them is empirical. GRU is **cheaper to run**; LSTM has **slightly more representational power**. **Try both and measure** rather than arguing about it.

Bidirectional layers: the same data, read backwards too

wrapping a layer

PYTHON

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Bidirectional(layers.LSTM(16))(inputs)
outputs = layers.Dense(1)(x)
```

On **this** problem it does not help: for weather, the recent past is what matters, so reading in reverse gives a **chronologically backwards view that is simply worse**. It shines on text, where both directions carry meaning — chapter 14 returns to it.

Going further

- ♦ **Adjust the number of units** in each recurrent layer, and the dropout rates — the configurations here were chosen largely arbitrarily.
- ♦ **Adjust the learning rate** of the optimizer.
- ♦ Try a **stack of `Dense` layers** as the regressor on top instead of a single one.
- ♦ **Improve the input**: try longer or shorter sequences, or a different sampling rate.

And the standing warning: **always evaluate on the validation set**, then confirm once on test — otherwise you end up **overfitting to your validation procedure**, exactly as chapter 5 described.

What has to survive this chapter

- 1 **Always establish a common-sense baseline first.** Here it beat two deep learning models.
- 2 **A good solution existing in your hypothesis space does not mean gradient descent will find it.**
- 3 **Dense flattens away time; 1D convolution pools away order.** Neither prior fits a forecasting problem.
- 4 **RNNs keep a state** and process a sequence incrementally.
- 5 **SimpleRNN cannot learn long dependencies;** LSTM adds a carry track — residual connections, along time.

...and the three refinements

Recurrent dropout

A **temporally constant mask**. A fresh random mask each timestep disrupts the error signal through time.

NOTEBOOK

[03_lstm_and_gru.ipynb](#)

Stacking

More capacity, at more compute. Intermediate layers need `return_sequences=True`.

NEXT

[Chapter 14 — Text classification](#)

Bidirectional

Helps where **both directions carry meaning** — text. Does not help where only the recent past matters.