



Institut Teknologi Bandung

Directorate of Sustainable Professional Education

CHAPTER 14

Text Classification

Tokenization, and the one question every NLP architecture answers differently: is a sentence a set of words, or a sequence of them?

Duration 3 hours (2 sessions)

Instructor Rahman Indra Kesuma, S.Kom., M.Cs. · Prof. Bambang Riyanto Trilaksono

Source Chollet & Watson, Deep Learning with Python 3e — chapter 14

Session Objectives

By the end of this session, participants can:

- 1 Explain why natural language is **messy in a way machine languages are not**, and why rule-based NLP failed.
- 2 Compare **character, word, and subword** tokenization, and say what byte-pair encoding optimises.
- 3 State the **set versus sequence** question and place bag-of-words, RNNs, and Transformers against it.
- 4 Build a **bag-of-words** classifier with `TextVectorization`, and extend it with **bigrams**.
- 5 Explain what a **word embedding** is, and why one-hot encoding makes a false assumption.
- 6 **Pretrain** an embedding with an unsupervised task (CBOW), and use it for classification.

BOOK

[Chapter 14 — full text](#)

NOTEBOOK

[01_tokenizers.ipynb](#)

NOTEBOOK

[02_bag_of_words_and_bigrams.ipynb](#)

NOTEBOOK

[03_sequence_models.ipynb](#)

NOTEBOOK

[04_pretraining_an_embedding.ipynb](#)

NOTEBOOK

[All chapter 14 notebooks](#)

REPO

[Author's official notebook](#)

Why "natural" language is a different kind of object

Machine languages were designed. An engineer wrote down formal rules first; people used the language once the rule set was complete.

Human language is the reverse. Usage comes first; rules arise later, are formalised after the fact, and are **often ignored or broken by its users**

So machine-readable language is highly structured and rigorous, while natural language is **messy — ambiguous, chaotic, sprawling, and constantly in flux.**

The 1954 demo, and the decade it bought


The goal was to attract excitement and funding, and in that sense **it was a huge success**. The authors claimed translation would be solved within five years. **Funding poured in for the better part of a decade.**

...and why it did not generalise

- ♦ Words change meaning dramatically depending on context.
- ♦ Every grammar rule needed countless exceptions.
- ♦ Shining on a few handpicked examples was simple; building a robust system that could compete with human translators **was another matter**.

An influential US report a decade later picked apart the lack of progress, and the funding dried up. Yet handcrafted rules **remained dominant well into the 1990s** — because there was no viable alternative.

The moment it became machine learning



Faster computers and more data in the late 1980s changed what was possible.

When you find yourself building systems that are big piles of ad hoc rules ... you're likely to start asking, "Could I use a corpus of data to automate the process of finding these rules?" And just like that, you've graduated to doing machine learning.

Chollet & Watson, section 14.1

01

Preparing text data

Tokenization: from characters to byte-pair encoding.

The simplest tokenizer, written out

listing 14.1 – a character-level tokenizer

PYTHON

```
class CharTokenizer:
    def __init__(self, vocabulary):
        self.vocabulary = vocabulary
        self.unk_id = vocabulary["[UNK]"]

    def standardize(self, inputs):
        return inputs.lower()

    def split(self, inputs):
        return re.findall(r".", inputs)

    def index(self, tokens):
        return [self.vocabulary.get(t, self.unk_id) for t in tokens]

    def __call__(self, inputs):
        return self.index(self.split(self.standardize(inputs)))
```

The `[UNK]` fallback is the important detail: **anything outside the vocabulary becomes one shared token**, and how often that happens is a property of the tokenizer's design.

Building a vocabulary, and capping it

computing a capped vocabulary

PYTHON

```
def compute_vocabulary(text_iterable, max_size, split_fn):
    counts = collections.Counter()
    for text in text_iterable:
        counts.update(split_fn(text.lower()))

    vocabulary = {"[UNK]": 0}  # index 0 is always the fallback
    for token, _ in counts.most_common(max_size - 1):
        vocabulary[token] = len(vocabulary)
    return vocabulary
```

Capping the vocabulary is the same **feature selection** move as chapter 5's `num_words=10000`: rare tokens carry little signal and **invite the spurious correlations chapter 5 warned about**

Word level: shorter sequences, a harder vocabulary problem

a word-level split


PYTHON

```
class WordTokenizer(CharTokenizer):
    def standardize(self, inputs):
        inputs = inputs.lower()
        return "".join(c for c in inputs if c not in string.punctuation)

    def split(self, inputs):
        return re.findall(r"\w+", inputs)
```

- ♦ **How large should the vocabulary be?** English has hundreds of thousands of word forms.
- ♦ **What happens to a word you have never seen?** It becomes [UNK], and **all its meaning is gone**.

Three levels, and the trade-off between them



Subword tokenization aims at the best of both.

Tokenization as compression

- ♦ **Shorter tokens** compress the overall length of each example.
- ♦ **A smaller vocabulary** reduces the bytes needed to represent each token.
- ♦ Achieve both and you feed the model **short, information-rich sequences**.

That analogy turned out to be powerful: one of the most effective tricks of the last decade was **repurposing a 1990s lossless compression algorithm — byte-pair encoding — for tokenization**. It is used by **ChatGPT and many other models to this day**

How byte-pair encoding builds its vocabulary

the BPE idea, in outline

PYTHON

```
# start: every character is a token
#   l o w e r   n e w e s t   w i d e s t

# merge the most frequent pair, repeatedly:
#   "e" + "s" -> "es"           n e w e s t   w i d e s t
#   "es" + "t" -> "est"        n e w e s t   w i d e s t
#   "l" + "o"  -> "lo"         l o w e r

# stop when the vocabulary reaches the target size
```

Common words end up as **single tokens**; rare words decompose into **familiar pieces**. Nothing is ever `[UNK]`, because **the character level is always available as a fallback**

02


Sets versus sequences

The pivotal question from which every NLP architecture springs.

Representing a token is easy. Representing order is not.

Unlike the steps of a timeseries, **words in a sentence have no natural, canonical order**. Different languages order similar words very differently; within one language you can usually say the same thing by reshuffling. **Order clearly matters, but its relationship to meaning is not straightforward.**

Three answers, three families of architecture



RNNs and Transformers both take order into account, so both are called sequence models.

Historically most early NLP was bag-of-words. Interest in sequence models only rose in **2015**, with the rebirth of RNNs.

Both approaches remain relevant today.

The benchmark: raw IMDB this time

listing 14.8 — downloading raw IMDB

PYTHON

```
zip_path = keras.utils.get_file(
    origin="https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz",
    fname="imdb",
    extract=True,
)
imdb_extract_dir = pathlib.Path(zip_path) / "aclImdb"

for path in imdb_extract_dir.glob("*/*"):
    if path.is_dir():
        print(path)
```

OUTPUT

```
aclImdb/train/pos
aclImdb/train/neg
aclImdb/train/unsup
aclImdb/test/pos
aclImdb/test/neg
```

Note `train/unsup` — 25,000 **unlabelled** reviews. They are ignored at first and become essential in section 14.5.4.

03

Set models

Throw order away, and see how far that gets you.

Bag of words, in one picture

The sentence

the cat sat on the mat

↓
count, and forget the order

The vector

1	1	1	1	2
cat	mat	on	sat	the

shape (5,) - one slot per vocabulary word, and this vocabulary has 5.

Reverse it and you cannot: cat×1 mat×1 on×1 sat×1 the×2 ... has no order left in it.

"the cat sat on the mat" and "the mat sat on the cat" give the identical vector. That is the whole limitation, in one line.

Tokenize, count, and let the order go.

The idea is simply to **assign a weight to every word**. *terrible* probably indicates a bad review; *riveting* probably indicates a good one.

TextVectorization does all of it

listing 14.12 — bag-of-words encoding

PYTHON

```
from keras import layers

max_tokens = 20_000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    split="whitespace",      # word-level vocabulary
    output_mode="multi_hot", # and multi-hot output, in one layer
)

train_ds_no_labels = train_ds.map(lambda x, y: x)
text_vectorization.adapt(train_ds_no_labels)
```

20,000 words is a good starting point for text classification. Note `adapt()` is given the reviews **without labels** — it is learning vocabulary, **not learning the task**

The model on top is tiny

the bag-of-words classifier


PYTHON

```
inputs = keras.Input(shape=(max_tokens,))
x = layers.Dense(16, activation="relu")(inputs)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)

model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])
```

It performs respectably — which is the point. **A model that cannot tell "good, not bad" from "bad, not good" still does well on sentiment**, because long reviews carry plenty of unordered evidence.

Bigrams: putting a little order back by hand



Each window shares 1 token(s) with the next. That overlap is what stops a boundary from cutting an idea in half — and what you pay for twice.

Every adjacent pair becomes a token of its own — and each pair shares a word with the next.

...and it is one argument in the layer

listing - bigram encoding

PYTHON

```
text_vectorization = layers.TextVectorization(  
    ngrams=2,                # unigrams AND bigrams  
    max_tokens=max_tokens,  
    output_mode="multi_hot",  
)
```

One argument, and the model improves: *"not good"* is now a token in its own right. But the approach **only scales to a local ordering of a few words**, and it is feature engineering done by hand.

Counts, or presence? A third output mode

output_mode	WHAT EACH POSITION HOLDS
"multi_hot"	1 if the word appears, 0 otherwise.
"count"	How many times it appears.
"tf_idf"	Its count, discounted by how common the word is across the whole corpus — so <i>the</i> contributes almost nothing and a distinctive word contributes a lot.

TF-IDF is decades older than deep learning and still competitive on short texts. It is worth trying before **anything with a recurrent layer in it**

04

Sequence models

Stop engineering order. Show the model the raw sequence.

The argument for moving on

As so often in deep learning: rather than building the features yourself, **expose the model to the raw word sequence and let it learn the positional dependencies** — **the same argument chapter 1 made about feature engineering in general**

One wrinkle: batches have to be rectangular

padding and truncation

PYTHON

```
max_length = 600

text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",          # integer token IDs, in order
    output_sequence_length=max_length, # pad or truncate to a fixed length
)
```

Padding introduces a new problem the model has to be told about: **those zeros are not words**. Getting this wrong means **the model spends capacity learning about padding**

A recurrent model on the raw sequence

a sequence model

PYTHON

```
inputs = keras.Input(shape=(max_length,), dtype="int32")
x = layers.Embedding(input_dim=max_tokens, output_dim=256, mask_zero=True)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
```

`mask_zero=True` is the answer to the padding problem: it tells every downstream layer to **skip the padded positions entirely**

What the sequence model buys, and what it costs

What it buys

Word order is **learned rather than engineered**, so dependencies longer than a bigram become reachable.

What it costs

Far slower to train than a bag-of-words model, and it needs **more data** to justify the extra parameters.

The book's standing advice applies: **start with the cheap model**. If bag-of-words is close to the sequence model on your data, **the sequence model is not earning its cost**

05

Word embeddings

One-hot encoding makes an assumption that is plainly false.

What one-hot encoding quietly assumes

For words that assumption is **clearly wrong**. "*Movie*" and "*film*" are interchangeable in most sentences, so their vectors **should be the same vector, or close to it** — not at right angles.

What an embedding is instead




Figure 14.3 — sparse, hardcoded, high-dimensional versus dense, learned, and lower-dimensional.

The goal: **the geometric relationship between two word vectors should reflect the semantic relationship between the words**. Related words close, unrelated words far apart.

Directions in the space mean things

- ♦ The same vector takes you from **cat** → **tiger** and from **dog** → **wolf** — a *from pet to wild animal* direction.
- ♦ Another takes you from **dog** → **cat** and **wolf** → **tiger** — a *from canine to feline* direction.
- ♦ In real embedding spaces the classic examples are **gender** and **plural** vectors: add *female* to *king* and you get **queen**.

Real word-embedding spaces typically feature **thousands** of such interpretable directions.

Is there one ideal embedding space?

The book's answer is careful: **possibly — but we have yet to compute anything of the sort.** What is available in practice is an embedding learned either **jointly with your task, or on a large corpus beforehand**

An `Embedding` layer trained with the classifier learns a space specialised to *this* task. That is often fine — and when labelled data is scarce, it is not.

Choosing the embedding dimension

DIMENSION	WHEN IT FITS
Too small	An information bottleneck — the space cannot separate words that mean different things. Chapter 4's 4-unit layer, again.
256 – 1,024	Typical for large vocabularies, and the range the book works in.
Too large	More parameters than the data can constrain; the embedding memorises rather than generalises.

For comparison: one-hot over this vocabulary would be **20,000 dimensions**. An embedding **packs more information into far fewer**

06

Pretraining

Where the unlabelled 25,000 reviews finally earn their place.

Why pretraining took over NLP

The idea: **devise an unsupervised task that needs no labels**. Pretraining data can be text from a similar domain, or even arbitrary text in the right language. It learns general patterns, **priming the model before you specialise it**

CBOW: guess the word from its neighbours




Figure 14.6 — Continuous Bag of Words, from Mikolov et al. (2013).

If the surrounding bag contains *sail*, *wave*, and *mast*, you might guess *boat*. **No labels are needed** — the supervision comes from the text itself, which is exactly chapter 1's **self-supervised learning**

Where the extra data comes from

combining labelled and unlabelled text

PYTHON

```
# all of train/pos, train/neg AND train/unsup, labels discarded
pretrain_ds = keras.utils.text_dataset_from_directory(
    imdb_extract_dir / "train",
    labels=None,
    batch_size=batch_size,
)
```

Doubling the text costs nothing, because the CBOW task does not care whether a review is positive or negative — **it only needs words in context**

Using the pretrained embedding

loading pretrained weights

PYTHON

```
embedding_layer = layers.Embedding(max_tokens, 256, mask_zero=True)
embedding_layer.build((None,))
embedding_layer.set_weights(pretrained_embedding.get_weights())

inputs = keras.Input(shape=(max_length,), dtype="int32")
x = embedding_layer(inputs)
x = layers.Bidirectional(layers.LSTM(32))(x)
outputs = layers.Dense(1, activation="sigmoid")(layers.Dropout(0.5)(x))
```

This is **transfer learning for text**, structurally identical to feature extraction from a pretrained ConvNet — **and it is the pattern the whole of chapter 15 builds on**

The four models, on one benchmark



Each step adds either more order information or more outside knowledge.

And the practical caveat: **the bag-of-words model is close behind, trains in seconds, and is trivial to explain.** On short texts, or with little data, **it is frequently the right answer**

Four ways text pipelines go wrong

Padding treated as content

Forgetting `mask_zero=True`. The model spends capacity modelling zeros, and long reviews are penalised against short ones.

`adapt()` on the wrong data

Calling `adapt()` on the full dataset builds the vocabulary from **test text** — an information leak in the sense of chapter 5.

Truncation that cuts the answer

`output_sequence_length` shorter than the reviews means the ending — often where sentiment resolves — is discarded.

A tokenizer trained on the wrong language

A vocabulary built on English text will decompose other languages into near-characters, inflating sequence length enormously.

What has to survive this chapter

- 1 **Natural language is messy by construction** — usage precedes rules, which is why rule-based NLP failed for forty years.
- 2 **Tokenization is compression.** Byte-pair encoding gets short sequences and a small vocabulary at once.
- 3 **Set or sequence** is the question every NLP architecture answers.
- 4 **Bag-of-words is a strong, cheap baseline.** Bigrams add local order for one argument.
- 5 **One-hot assumes words are independent**, which is false. Embeddings learn a space where geometry reflects meaning.
- 6 **Pretraining on unlabelled text** removes the labelled-data bottleneck.

NOTEBOOK

[04_pretraining_an_embedding.ipynb](#)

NEXT

[Chapter 15 — Language models and the Transformer](#)